The Importance of Communities for Learning to Influence

Eric Balkanski* Nicole Immorlica[†] Yaron Singer[‡]

Abstract

We consider the canonical problem of influence maximization in social networks. Since the seminal work of Kempe, Kleinberg, and Tardos, there have been two largely disjoint efforts on this problem. The first studies the problem associated with learning the parameters of the generative influence model. The second focuses on the algorithmic challenge of identifying a set of influencers, assuming the parameters of the generative model are known. Recent results on learning and optimization imply that in general, if the generative model is not known but rather learned from training data, no algorithm can yield a constant factor approximation guarantee using polynomially-many samples, drawn from any distribution.

In this paper, we design a simple heuristic that overcomes this negative result in practice by leveraging the strong community structure of social networks. Although in general the approximation guarantee of our algorithm is necessarily unbounded, we show that this algorithm performs well experimentally. To justify its performance, we prove our algorithm obtains a constant factor approximation guarantee on graphs generated through the stochastic block model, traditionally used to model networks with community structure.

^{*}Harvard University, ericbalkanski@g.harvard.edu.

[†]Microsoft Research, nicimm@microsoft.com.

[‡]Harvard University, yaron@seas.harvard.edu.

1 Introduction

For well over a decade now, there has been extensive work on the canonical problem of influence maximization in social networks. First posed by Domingos and Richardson [DR01, RD02] and elegantly formulated and further developed by Kempe, Kleinberg, and Tardos [KKT03], influence maximization is the algorithmic challenge of selecting individuals who can serve as early adopters of a new idea, product, or technology in a manner that will trigger a large cascade in the social network.

In their seminal paper, Kempe, Kleinberg, and Tardos characterize a family of natural influence processes for which selecting a set of individuals that maximize the resulting cascade reduces to maximizing a submodular function under a cardinality constraint. Since submodular functions can be maximized within a 1-1/e approximation guarantee, one can then obtain desirable guarantees for the influence maximization problem. There have since been two, largely separate, agendas of research on the problem. The first line of work is concerned with learning the underlying submodular function from observations of cascades [LK03, AA05, LMF⁺07, GBL10, CKL11, GBS11, NS12, GLK12, DSSY12, ACKP13, DSGRZ13, FK14, DBB⁺14, CAD⁺14, DGSS14, DLBS14, NPS15, HO15]. The second line of work focuses on algorithmic challenges revolving around maximizing influence, assuming the underlying function that generates the diffusion process is known [KKT05, MR07, SS13, BBCL14, HS15, HK16, AS16].

In this paper, we consider the problem of learning to influence where the goal is to maximize influence from observations of cascades. This problem synthesizes both problems of learning the function from training data and of maximizing influence given the influence function. A natural approach for learning to influence is to first learn the influence function from cascades, and then apply a submodular optimization algorithm on the function learned from data. Somewhat counter-intuitively, it turns out that this approach yields desirable guarantees only under very strong learnability conditions¹. In some cases, when there are sufficiently many samples, and one can observe exactly which node attempts to influence whom at every time step, these learnability conditions can be met. A slight relaxation however (e.g. when there are only partial observations [NPS15, HXKL16]), can lead to sharp inapproximability.

A recent line of work shows that even when a function is statistically learnable, optimizing the function learned from data can be inapproximable [BRS17, BS17]. In particular, even when the submodular function $f: 2^N \to \mathbb{R}$ is a coverage function (which is PMAC learnable [BDF⁺12, FK14]), one would need to observe exponentially many samples $\{S_i, f(S_i)\}_{i=1}^m$ to obtain a constant factor approximation guarantee. Since coverage functions are special cases of the well studied models of influence (independent cascade, linear and submodular threshold), this implies that when the influence function is not known but learned from data, the influence maximization problem is intractable.

Learning to influence social networks. As with all impossibility results, the inapproximability discussed above holds for worst case instances, and it may be possible that such instances are rare for influence in social networks. In recent work, it was shown that when a submodular function has bounded curvature, there is a simple algorithm that can maximize the function under a cardinality

¹In general, the submodular function $f: 2^N \to \mathbb{R}$ needs to be learnable *everywhere* within arbitrary precision, i.e. for every set S one needs to assume that the learner can produce a surrogate function $\tilde{f}: 2^N \to \mathbb{R}$ s.t. for every $S \subseteq N$ the surrogate guarantees to be $(1 - \epsilon)f(S) \le \tilde{f}(S) \le (1 + \epsilon)f(S)$, for $\epsilon \in o(1)$ [HS16, HS17].

constraint from samples [BRS16]. Unfortunately, simple examples show that submodular functions that dictate influence processes in social networks do not have bounded curvature. Are there other reasonable conditions on social networks that yield desirable approximation guarantees?

Main result. In this paper we present a simple algorithm for learning to influence. This algorithm leverages the idea that social networks exhibit strong community structure. At a high level, the algorithm observes cascades and aims to select a set of nodes that are influential, but belong to different communities. Intuitively, when an influential node from a certain community is selected to initiate a cascade, the marginal contribution of adding another node from that same community is small, since the nodes in that community were likely already influenced. This observation can be translated into a simple algorithm which performs very well in practice. Analytically, since community structure is often modeled using stochastic block models, we prove that the algorithm obtains a constant factor approximation guarantee in such models, under mild assumptions.

1.1 Technical overview

The analysis for the approximation guarantees lies at the intersection of combinatorial optimization and random graph theory. We formalize the intuition that the algorithm leverages the community structure of social networks in the standard model to analyze communities, which is the stochastic block model. Intuitively, the algorithm obtains good approximations by picking the nodes that have the largest individual influence while avoiding picking multiple nodes in the same community by pruning nodes with high influence overlap. The individual influence of nodes and their overlap are estimated by the algorithm with what we call first and second order marginal contributions of nodes, which can be estimated from samples. We then uses phase transition results of Erdős–Rényi random graphs and branching processes techniques to compare these individual influences for nodes in different communities in the stochastic block model and bound the overlap of pairs of nodes.

The optimization from samples model. Optimization from samples was recently introduced by [BRS17] in the context of submodular optimization, we give the definition for general set functions.

Definition 1. A class of functions $\mathcal{F} = \{f : 2^N \to \mathbb{R}\}$ is α -optimizable from samples over distribution \mathcal{D} under constraint \mathcal{M} if there exists an algorithm s.t. for all $f \in \mathcal{F}$, given a set of samples $\{(S_i, f(S_i))\}_{i=1}^m$ where the sets S_i are drawn i.i.d. from \mathcal{D} , the algorithm returns $S \in \mathcal{M}$ s.t.:

$$\Pr_{S_1, \dots, S_m \sim \mathcal{D}} \left[\mathbb{E}[f(S)] \ge \alpha \cdot \max_{T \in \mathcal{M}} f(T) \right] \ge 1 - \delta,$$

where the expectation is over the decisions of the algorithm and $m \in \text{poly}(|N|, 1/\delta)$.

We focus on bounded product distributions \mathcal{D} , so every node a is, independently, in $S \sim \mathcal{D}$ with some probability $p_a \in [1/\operatorname{poly}(n), 1-1/\operatorname{poly}(n)]$. We assume this is the case throughout the paper.

Influence process. We assume that the influence process follows the standard independent cascade model. In the independent cascade model, a node a influences each of its neighbors b with some probability q_{ab} , independently. Thus, given a seed set of nodes S, the set of nodes influenced is the number of nodes connected to some node in S in the random subgraph of the network which contains every edge ab independently with probability q_{ab} . We define f(S) to be the expected number of nodes influenced by S according to the independent cascade model over some weighted social network.

The learning to influence model: optimization from samples for influence maximization.

The learning to influence model is an interpretation of the optimization from samples model [BRS17] for the specific problem of influence maximization in social networks. We are given a collection of samples $\{(S_i, |\operatorname{cc}(S_i)|)\}_{i=1}^m$ where sets S_i are the seed sets of nodes and $|\operatorname{cc}(S_i)|$ is the number of nodes influenced by S_i , i.e., the number of nodes that are connected to S_i in the random subgraph of the network. This number of nodes is a random variable with expected value $f(S_i) := \mathbb{E}[|\operatorname{cc}(S_i)|]$ over the realization of the influence process. Each sample is an independent realization of the influence process. The goal is then to find a set of nodes S under a cardinality constraint k which maximizes the influence in expectation, i.e., find a set S of size at most k which maximizes the expected number of nodes f(S) influenced by seed set S.

2 The Algorithm

We present the main algorithm, COPS. This algorithm is based on a novel optimization from samples technique which detects overlap in the marginal contributions of two different nodes, which is useful to avoid picking two nodes who have intersecting influence over a same collection of nodes.

2.1 Description of COPS

COPS, consists of two steps. It first orders nodes in decreasing order of first order marginal contribution, which is the expected marginal contribution of a node a to a random set $S \sim \mathcal{D}$. Then, it iteratively removes nodes a whose marginal contribution overlaps with the marginal contribution of at least one node before a in the ordering. The solution is the k first nodes in the pruned ordering.

Algorithm 1 COPS, learns to influence networks with COmmunity Pruning from Samples.

Input: Samples $S = \{(S, f(S))\}$, acceptable overlap α .

Order nodes according to their first order marginal contributions

Iteratively remove from this ordering nodes a whose marginal contribution has overlap of at least α with at least one node before a in this ordering.

return k first nodes in the ordering

The strong performance of this algorithm for the problem of influence maximization is best explained with the concept of communities. Intuitively, this algorithm first orders nodes in decreasing order of their individual influence and then removes nodes which are in a same community. This second step allows the algorithm to obtain a diverse solution which influences multiple different communities of the social network. In comparison, previous algorithms in optimization from samples [BRS16, BRS17] only use first order marginal contributions and perform well if the function is close to linear. Due to the high overlap in influence between nodes in a same community, influence functions are far from being linear and these algorithms have poor performance for influence maximization since they only pick nodes from a very small number of communities.

2.2 Computing overlap using second order marginal contributions

We define second order marginal contributions, which are used to compute the overlap between the marginal contribution of two nodes.

Definition 2. The second order expected marginal contribution of a node a to a random set S containing node b is

$$v_b(a) := \underset{S \sim \mathcal{D}: a \notin S, b \in S}{\mathbb{E}} [f(S \cup \{a\}) - f(S)].$$

The first order marginal contribution v(a) of node a is defined similarly as the marginal contribution of a node a to a random set S, i.e., $v(a) := \mathbb{E}_{S \sim \mathcal{D}: a \notin S}[f(S \cup \{a\}) - f(S)]$. These contributions can be estimated arbitrarily well for product distributions \mathcal{D} by taking the difference between the average value of samples containing a and b and the average value of samples containing b but not a (see Appendix B for details).

The subroutine OVERLAP (a, b, α) , $\alpha \in [0, 1]$, compares the second order marginal contribution of a to a random set containing b and the first order marginal contribution of a to a random set. If b causes the marginal contribution of a to decrease by at least a factor of $1 - \alpha$, then we say that a has marginal contribution with overlap of at least α with node b.

Algorithm 2 OVERLAP (a, b, α) , returns true if a and b have marginal contributions that overlap by at least a factor α .

Input: Samples $S = \{(S, f(S))\}$, node a, acceptable overlap α

If second order marginal contribution $v_b(a)$ is at least a factor of $1 - \alpha$ smaller than first order marginal contribution v(a),

return Node a has overlap of at least α with node b

OVERLAP is used to detect nodes in a same community. In the extreme case where two nodes a and b are in a community C where any node in C influences all of community C, then the second order marginal contribution $v_b(a)$ of a to random set S containing b is $v_b(a) = 0$ since b already influences all of C so a does not add any value, while $v(a) \approx |C|$. In the opposite case where a and b are in two communities which are not connected in the network, we have $v(a) = v_b(a)$ since adding b to a random set S has no impact on the value added by a.

2.3 Analyzing community structure

The main benefit from COPS is that it leverages the community structure of social networks. To formalize this explanation, we analyze our algorithm in the standard model used to study the community structure of networks, the stochastic block model. In this model, a fixed set of nodes V is partitioned in communities C_1, \ldots, C_ℓ . The network is then a random graph G = (V, E) where edges are added to E independently and where an intra-community edge is in E with much larger probability than an inter-community edge. These edges are added with identical probability q_C^{sb} for every edge in a same community, but with different probabilities for edges inside different communities C_i and C_j . We illustrate this model in Figure 1.

3 Dense Communities and Small Seed Set in the Stochastic Block Model

In this section, we show that COPS achieves a $1 - O(|C_k|^{-1})$ approximation, where C_k is the kth largest community, in the regime with dense communities and small seed set, which is described below. We show that the algorithm picks a node from each of the k largest communities with

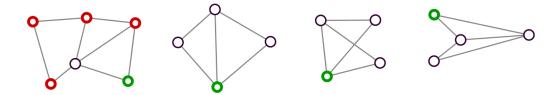


Figure 1: An illustration of the stochastic block model with communities C_1 , C_2 , C_3 and C_4 of sizes 6, 4, 4 and 4. The optimal solution for influence maximization with k = 4 is in green. Picking the k first nodes in the ordering by marginal contributions without pruning, as in [BRS16], leads to a solution with nodes from only C_1 (red). By removing nodes with overlapping marginal contributions, COPS obtains a diverse solution.

high probability, which is the optimal solution. In the next section, we show a constant factor approximation algorithm for a generalization of this setting, which requires a more intricate analysis.

In order to focus on the main characteristics of the community structure as an explanation for the performance of the algorithm, we make the following simplifying assumptions for the analysis. We first assume that there are no inter-community edges. We also assume that the random graph obtained from the stochastic block model is redrawn for every sample and that we aim to find a good solution in expectation over both the stochastic block model and the independent cascade model.

Formally, let G = (V, E) be the random graph over n nodes obtained from an independent cascade process over the graph generated by the stochastic block model. Similarly as for the stochastic block model, edge probabilities for the independent cascade model may vary between different communities and are identical within a single community C, where all edges have weights $q_C^{\rm ic}$. Thus, an edge e between two nodes in a community C is in E with probability $p_C := q_C^{\rm ic} \cdot q_C^{\rm sb}$, independently for every edge, where $q_C^{\rm ic}$ and $q_C^{\rm sb}$ are the edge probabilities in the independent cascade model and the stochastic block model respectively. The total influence by seed set S is then $|cc_G(S_i)|$ where $cc_G(S)$ is the set of nodes connected to S in G and we drop the subscript when it is clear from context. Thus, the objective function is $f(S) := \mathbb{E}_G[|cc(S)|]$. We describe the two assumptions for this section.

Dense communities. We assume that for the k largest communities C, $p_C > 3 \log |C|/|C|$ and C has super-constant size ($|C| = \omega(1)$). This assumption corresponds to communities where the probability p_C that a node $a_i \in C$ influences another node $a_j \in C$ is large. Since the subgraph G[C] of G induced by a community C is an Erdős–Rényi random graph, we get that G[C] is connected with high probability (see Appendix C).

Lemma 3. [ER60] Assume C is a "dense" community, then the subgraph G[C] of G is connected with probability $1 - O(|C|^{-2})$.

Small seed set. We also assume that the seed sets $S \sim \mathcal{D}$ are small enough so that they rarely intersect with a fixed community C, i.e., $\Pr_{S \sim \mathcal{D}}[S \cap C = \emptyset] \ge 1 - o(1)$. This assumption corresponds to cases where the set of early influencers is small, which is usually the case in cascades.

The analysis in this section relies on two main lemmas. We first show that the first order marginal contribution of a node is approximately the size of the community it belongs to (Lemma 4). Thus,

²The analysis easily extends to cases where inter-community edges form with probability significantly smaller to $q_C^{\rm sb}$, for all C.

the ordering by marginal contributions orders elements by the size of the community they belong to. Then, we show that any node $a \in C$ that is s.t. that there is a node $b \in C$ before a in the ordering is pruned (Lemma 5). Regarding the distribution $S \sim \mathcal{D}$ generating the samples, as previously mentioned, we consider any bounded product distribution. This implies that w.p. $1 - 1/\operatorname{poly}(n)$, the algorithm can compute marginal contribution estimates \tilde{v} that are all a $1/\operatorname{poly}(n)$ -additive approximation to the true marginal contributions v (See Appendix B for formal analysis of estimates). Thus, we give the analysis for the true marginal contributions, which, with probability $1 - 1/\operatorname{poly}(n)$ over the samples, easily extends for arbitrarily good estimates.

The following lemma shows that the ordering by first order marginal contributions corresponds to the ordering by decreasing order of community sizes that nodes belong to.

Lemma 4. For all $a \in C$ where C is one of the k largest communities, the first order marginal contribution of node a is approximately the size of its community, i.e., $(1 - o(1))|C| \le v(a) \le |C|$.

Proof. Assume a is a node in one of the k largest communities. Let \mathcal{D}_a and \mathcal{D}_{-a} denote the distributions $S \sim \mathcal{D}$ conditioned on $a \in S$ and $a \notin S$ respectively. We also denote marginal contributions by $f_S(a) := f(S \cup \{a\}) - f(S)$. We obtain

$$v(a) = \underset{S \sim \mathcal{D}_{-a}, G}{\mathbb{E}}[f_S(a)] \ge \underset{S \sim \mathcal{D}_{-a}}{\Pr}[S \cap C = \emptyset] \cdot \underset{G}{\Pr}[\operatorname{cc}(a) = C] \cdot \underset{S \sim \mathcal{D}_{-a}: S \cap C = \emptyset, \\ G : \operatorname{cc}(a) = C}{\mathbb{E}}[f_S(a)]$$
$$= \underset{S \sim \mathcal{D}_{-a}}{\Pr}[S \cap C = \emptyset] \cdot \underset{G}{\Pr}[\operatorname{cc}(a) = C] \cdot |C|$$
$$\ge (1 - o(1)) \cdot |C|$$

where the last inequality is by the small seed set assumption and since C is connected with probability 1 - o(1) (Lemma 3 and $|C| = \omega(1)$ by dense community assumption). For the upper bound, v(a) is trivially at most the size of a's community since there are no inter-community edges.

The next lemma shows that the algorithm does not pick two nodes in a same community.

Lemma 5. With probability 1 - o(1), for all pairs of nodes a, b such that $a, b \in C$ where C is one of the k largest communities, OVERLAP $(a, b, \alpha) = True$ for any constant $\alpha \in [0, 1)$.

Proof. Let a, b be two nodes in one of the k largest communities C and $\mathcal{D}_{-a,b}$ denote the distribution $S \sim \mathcal{D}$ conditioned on $a \notin S$ and $b \in S$. Then,

$$v_b(a) = \underset{S \sim \mathcal{D}_{-a,b}}{\mathbb{E}} [f_S(a)] \le \Pr[b \in \operatorname{cc}(a)] \cdot 0 + \Pr[b \notin \operatorname{cc}(a)] \cdot |C| = o(1) \le o(1) \cdot v(a)$$

where the last equality is since G[C] is not connected w.p. $O(|C|^{-2})$ by Lemma 3 and since $|C| = \omega(1)$ by the dense community assumption, which concludes the proof.

By combining Lemmas 4 and 5, we obtain the main result for this section (proof in Appendix D).

Theorem 6. In the dense communities and small seed set setting, COPS with α -overlap allowed, for any constant $\alpha \in (0,1)$ is a 1-o(1)-approximation algorithm for learning to influence from samples from a bounded product distribution \mathcal{D} .

4 Constant Approximation for General Stochastic Block Model

In this section, we relax assumptions from the previous section and show that COPS is a constant factor approximation algorithm in this more demanding setting. Recall that G is the random graph obtained from both the stochastic block model and the independent cascade model. A main observation that is used in the analysis is to observe that the random subgraph G[C], for some community C, is an Erdős–Rényi random graph $G_{|C|,p_C}$.

Relaxation of the assumptions. Instead of only considering dense communities where $p_C = \Omega((\log |C|)/|C|)$, we consider both tight communities C where $p_C \geq (1+\epsilon)/|C|$ for some constant $\epsilon > 0$ and loose communities C where $p_C \leq (1-\epsilon)/|C|$ for some constant $\epsilon > 0$. We also relax the small seed set assumption to the reasonable non-ubiquitous seed set assumption. Instead of having a seed set $S \sim \mathcal{D}$ rarely intersect with a fixed community C, we only assume that $\Pr_{S \sim \mathcal{D}}[S \cap C = \emptyset] \geq \epsilon$ for some constant $\epsilon > 0$. Again, since seed sets are of small sizes in practice, it seems reasonable that with some constant probability a community does not contain any seeds.

Overview of analysis. At a high level, the analysis exploits the remarkably sharp threshold for the phase transition of Erdős-Rényi random graphs. This phase transition (Lemma 7) tells us that a tight community C contains w.h.p. a giant connected component with a constant fraction of the nodes from C. Thus, a single node from a tight community influences a constant fraction of its community in expectation. The ordering by first order marginal contributions thus ensures a constant factor approximation of the value from nodes in tight communities (Lemma 10). On the other hand, we show that a node from a loose community influences only at most a constant number of nodes in expectation (Lemma 8) by using branching processes. Since the algorithm checks for overlap using second order marginal contributions, the algorithm picks at most one node from any tight community (Lemma 11). Combining all the pieces together, we obtain a constant factor approximation (Theorem 12).

We first state the result for the giant connected component in a tight community, which is an immediate corollary of the prominent giant connected component result in the Erdős–Rényi model.

Lemma 7. [ER60] Let C be a tight community with $|C| = \omega(1)$, then G[C] has a "giant" connected component containing a constant fraction of the nodes in C w.p. 1 - o(1).

The following lemma analyzes the influence of a node in a loose community through the lenses of Galton-Watson branching processes to show that such a node influences at most a constant number of nodes in expectation. The proof is deferred to Appendix E.

Lemma 8. Let C be a loose community, then $f(\{a\}) \leq c$ for all $a \in C$ and some constant c.

We can now upper bound the value of the optimal solution S^* . Let C_1, \ldots, C_t be the $t \leq k$ tight communities that have at least one node in C_i that is in the optimal solution S^* and that are of super-constant size, i.e., $|C| = \omega(1)$. Without loss, we order these communities in decreasing order of their size $|C_i|$.

Lemma 9. Let S^* be the optimal set of nodes and C_i and t be defined as above. There exists a constant c such that $f(S^*) \leq \sum_{i=1}^t |C_i| + c \cdot k$.

Thus, we consider all possible sizes of communities except communities of size that converges to exactly $1/p_C$, which is unlikely to occur in practice.

Proof. Let S_A^{\star} and S_B^{\star} be a partition of the optimal nodes in nodes that are in tight communities with super-constant individual influence and nodes that are not in such a community. The influence $f(S_A^{\star})$ is trivially upper bounded by $\sum_{i=1}^t |C_i|$. Next, there exists some constant c s.t. $f(S_B^{\star}) \leq \sum_{a \in S_B^{\star}} f(\{a\}) \leq c$ where the first inequality is by submodularity and the second since nodes in loose communities have constant individual influence by Lemma 8 and nodes in tight community without super-constant individual influence have constant influence by definition. We conclude that by submodularity, $f(S^{\star}) \leq f(S_A^{\star}) + f(S_B^{\star}) \leq \sum_{i=1}^t |C_i| + c \cdot k$.

Next, we argue that the solution returned by the algorithm is a constant factor away from $\sum_{i=1}^{t} |C_i|$.

Lemma 10. Let a be the ith node in the ordering by first order maginal contribution after the pruning and C_i be the ith largest tight community with super-constant individual influence and with at least one node in the optimal solution S^* . Then, $f(\{a\}) \ge \epsilon |C_i|$ for some constant $\epsilon > 0$.

Proof. By definition of C_i , we have $|C_1| \ge \cdots \ge |C_i|$ that are all tight communities. Let b be a node in C_j for $j \in [i]$, $\mathbb{1}_{gc(C)}$ be the indicator variable indicating if there is a giant component in community C, and gc(C) be this giant component. We get

$$v(b) \ge \Pr[\mathbb{1}_{\gcd(C_j)}] \cdot \Pr_{S \sim \mathcal{D}_{-b}}[S \cap C_j = \emptyset] \cdot \Pr[b \in \gcd(C_j)] \cdot \mathbb{E}[|\gcd(C_j)| : b \in \gcd(C_j)]$$

$$\ge (1 - o(1)) \cdot \epsilon_1 \cdot \epsilon_2 \cdot \epsilon_3 |C_j| \ge \epsilon |C_j|$$

for some constants $\epsilon_1, \epsilon_2, \epsilon_3, \epsilon > 0$ by Lemma 7 and the non-ubiquitous assumption. Similarly as in Theorem 6, if a and b are in different communities, OVERLAP $(a, b, \alpha) = \text{False}$ for $\alpha \in (0, 1]$. Thus, there is at least one node $b \in \bigcup_{j=1}^{i} C_j$ at position i or after in the ordering after the pruning, and $v(b) \ge \epsilon |C_j|$ for some $j \in [i]$. By the ordering by first order marginal contributions and since node a is in ith position, $v(a) \ge v(b)$, and we get that $f(\{a\}) \ge v(a) \ge v(b) \ge \epsilon |C_i| \ge \epsilon |C_i|$.

Next, we show that the algorithm never picks two nodes from a same tight community and defer the proof to Appendix E.

Lemma 11. If $a, b \in C$ and C is a tight community, then OVERLAP $(a, b, \alpha) = True$ for $\alpha = o(1)$.

We combine the above lemmas to obtain the approximation guarantee of COPS (proof in Appendix E).

Theorem 12. With overlap allowed $\alpha = 1/\operatorname{poly}(n)$, COPS is a constant factor approximation algorithm for learning to influence from samples drawn from a bounded product distribution \mathcal{D} in the setting with tight and loose communities and non-ubiquitous seed sets.

5 Experiments

In this section, we compare the performance of COPS and three other algorithms on real and synthetic networks. We show that COPS performs well in practice, it outperforms the previous optimization from samples algorithm and gets closer to the solution obtained when given complete access to the influence function.

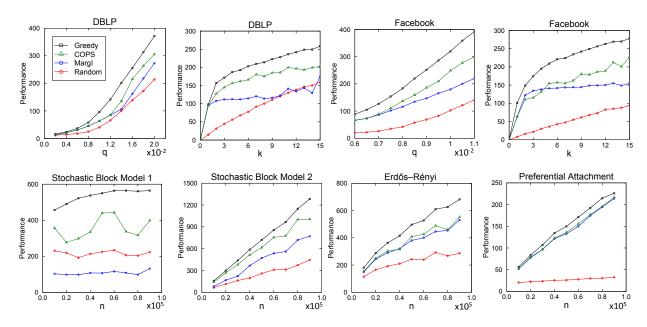


Figure 2: Empirical performance of COPS against the GREEDY upper bound, the previous optimization from samples algorithm MARGI and a random set.

Experimental setup. The first synthetic network considered is the stochastic block model, SBM 1, where communities have random sizes with one community of size significantly larger than the other communities. We maintained the same expected community size as n varied. In the second stochastic block model, SBM 2, all communities have same expected size and the number of communities was fixed as n varied. The third and fourth synthetic networks were an Erdős–Rényi (ER) random graph and the preferential attachment model (PA). Experiments were also conducted on two real networks publicly available ([LK15]). The first is a subgraph of the Facebook social network with n = 4k and m = 88k. The second is a subgraph of the DBLP co-authorship network, which has ground truth communities as described in [LK15], where nodes of degree at most 10 were pruned to obtain n = 54k, m = 361k and where the 1.2k nodes with degree at least 50 were considered as potential nodes in the solution.

Benchmarks. We considered three different benchmarks to compare the COPS algorithm against. The standard Greedy algorithm in the value query model is an upper bound since it is the optimal efficient algorithm given value query access to the function and COPS is in the more restricted setting with only samples. MargI is the optimization from samples algorithm which picks the k nodes with highest first order marginal contribution ([BRS16]) and does not use second order marginal contributions. Random simply returns a random set. All the samples are drawn from the product distribution with marginal probability k/n, so that samples have expected size k. We further describe the parameters of each plot in Appendix F.

Empirical evaluation. COPS significantly outperforms the previous optimization from samples algorithm MARGI, getting much closer to the GREEDY upper bound. We observe that the more there is a community structure in the network, the better the performance of COPS is compared to MARGI, e.g., SBM vs ER and PA (which do not have a community structure). When the edge

weight $q:=q^{\text{i.c.}}$ for the cascades is small, the function is near-linear and MARGI performs well, whereas when it is large, there is a lot of overlap and COPS performs better. The performance of COPS as a function of the overlap allowed (experiment in Appendix F) can be explained as follows: Its performance slowly increases as the the overlap allowed increases and COPS can pick from a larger collection of nodes until it drops when it allows too much overlap and picks mostly very close nodes from a same community. For SBM 1 with one larger community, MARGI is trapped into only picking nodes from that larger community and performs even less well than RANDOM. As n increases, the number of nodes influenced increases roughly linearly for SBM 2 when the number of communities is fixed since the number of nodes per community increases linearly, which is not the case for SBM 1.

References

- [AA05] Eytan Adar and Lada A. Adamic. Tracking information epidemics in blogspace. In WI, 2005.
- [ACKP13] Bruno D. Abrahao, Flavio Chierichetti, Robert Kleinberg, and Alessandro Panconesi. Trace complexity of network inference. In KDD, 2013.
- [AS16] Rico Angell and Grant Schoenebeck. Don't be greedy: Leveraging community structure to find high quality seed sets for influence maximization. arXiv preprint arXiv:1609.06520, 2016.
- [BBCL14] Christian Borgs, Michael Brautbar, Jennifer T. Chayes, and Brendan Lucier. Maximizing social influence in nearly optimal time. In *Proceedings of the Twenty-Fifth Annual ACM-SIAM Symposium on Discrete Algorithms*, SODA 2014, Portland, Oregon, USA, January 5-7, 2014, pages 946–957, 2014.
- [BDF⁺12] Ashwinkumar Badanidiyuru, Shahar Dobzinski, Hu Fu, Robert Kleinberg, Noam Nisan, and Tim Roughgarden. Sketching valuation functions. In *Proceedings of the twenty-third annual ACM-SIAM symposium on Discrete Algorithms*, pages 1025–1035. Society for Industrial and Applied Mathematics, 2012.
- [BHK] Avrim Blum, John Hopcroft, and Ravindran Kannan. Foundations of data science.
- [BRS16] Eric Balkanski, Aviad Rubinstein, and Yaron Singer. The power of optimization from samples. In Advances in Neural Information Processing Systems 29: Annual Conference on Neural Information Processing Systems 2016, December 5-10, 2016, Barcelona, Spain, pages 4017–4025, 2016.
- [BRS17] Eric Balkanski, Aviad Rubinstein, and Yaron Singer. The limitations of optimization from samples. In STOC, 2017.
- [BS17] Eric Balkanski and Yaron Singer. The sample complexity of optimizing a convex function. In *COLT*, 2017.
- [CAD+14] Justin Cheng, Lada A. Adamic, P. Alex Dow, Jon M. Kleinberg, and Jure Leskovec. Can cascades be predicted? In WWW, 2014.
- [CKL11] Flavio Chierichetti, Jon M. Kleinberg, and David Liben-Nowell. Reconstructing patterns of information diffusion from incomplete observations. In NIPS, 2011.
- [DBB⁺14] Abir De, Sourangshu Bhattacharya, Parantapa Bhattacharya, Niloy Ganguly, and Soumen Chakrabarti. Learning a linear influence model from transient opinion dynamics. In CIKM, 2014.
- [DGSS14] Hadi Daneshmand, Manuel Gomez-Rodriguez, Le Song, and Bernhard Schölkopf. Estimating diffusion network structures: Recovery conditions, sample complexity & soft-thresholding algorithm. In ICML, 2014.
- [DLBS14] Nan Du, Yingyu Liang, Maria-Florina Balcan, and Le Song. Influence function learning in information diffusion networks. In ICML, 2014.
- [DR01] Pedro Domingos and Matthew Richardson. Mining the network value of customers. In KDD, 2001.
- [DSGRZ13] Nan Du, Le Song, Manuel Gomez-Rodriguez, and Hongyuan Zha. Scalable influence estimation in continuous-time diffusion networks. In NIPS, 2013.
- [DSSY12] Nan Du, Le Song, Alexander J. Smola, and Ming Yuan. Learning networks of heterogeneous influence. In NIPS, 2012.
- [ER60] Paul Erdos and Alfréd Rényi. On the evolution of random graphs. *Publ. Math. Inst. Hung. Acad. Sci*, 5(1):17–60, 1960.
- [FK14] Vitaly Feldman and Pravesh Kothari. Learning coverage functions and private release of marginals. In COLT, 2014.
- [GBL10] Amit Goyal, Francesco Bonchi, and Laks VS Lakshmanan. Learning influence probabilities in social networks. In KDD, 2010.
- [GBS11] Manuel Gomez-Rodriguez, David Balduzzi, and Bernhard Schölkopf. Uncovering the temporal dynamics of diffusion networks. In *ICML*, 2011.
- [GLK12] Manuel Gomez-Rodriguez, Jure Leskovec, and Andreas Krause. Inferring networks of diffusion and influence. ACM Transactions on Knowledge Discovery from Data, 5(4):21, 2012.

- [HK16] Xinran He and David Kempe. Robust influence maximization. In Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, San Francisco, CA, USA, August 13-17, 2016, pages 885–894, 2016.
- [HO15] Jean Honorio and Luis Ortiz. Learning the structure and parameters of large-population graphical games from behavioral data. *Journal of Machine Learning Research*, 16:1157–1210, 2015.
- [HS15] Thibaut Horel and Yaron Singer. Scalable methods for adaptively seeding a social network. In *Proceedings* of the 24th International Conference on World Wide Web, WWW 2015, Florence, Italy, May 18-22, 2015, pages 441-451, 2015.
- [HS16] Thibaut Horel and Yaron Singer. Maximization of approximately submodular functions. In Advances in Neural Information Processing Systems 29: Annual Conference on Neural Information Processing Systems 2016, December 5-10, 2016, Barcelona, Spain, pages 3045–3053, 2016.
- [HS17] Avinatan Hassidim and Yaron Singer. Submodular maximization under noise. In COLT, 2017.
- [HXKL16] Xinran He, Ke Xu, David Kempe, and Yan Liu. Learning influence functions from incomplete observations. In Advances in Neural Information Processing Systems 29: Annual Conference on Neural Information Processing Systems 2016, December 5-10, 2016, Barcelona, Spain, pages 2065–2073, 2016.
- [KKT03] David Kempe, Jon M. Kleinberg, and Éva Tardos. Maximizing the spread of influence through a social network. In KDD, 2003.
- [KKT05] David Kempe, Jon M. Kleinberg, and Éva Tardos. Influential nodes in a diffusion model for social networks. In Automata, Languages and Programming, 32nd International Colloquium, ICALP 2005, Lisbon, Portugal, July 11-15, 2005, Proceedings, pages 1127–1138, 2005.
- [LK03] David Liben-Nowell and Jon M. Kleinberg. The link prediction problem for social networks. In CIKM, 2003.
- [LK15] Jure Leskovec and Andrej Krevl. Snap datasets, stanford large network dataset collection. 2015.
- [LMF⁺07] Jure Leskovec, Mary McGlohon, Christos Faloutsos, Natalie S. Glance, and Matthew Hurst. Patterns of cascading behavior in large blog graphs. In SDM, 2007.
- [MR07] Elchanan Mossel and Sébastien Roch. On the submodularity of influence in social networks. In STOC, 2007.
- [NPS15] Harikrishna Narasimhan, David C. Parkes, and Yaron Singer. Learnability of influence in networks. In Advances in Neural Information Processing Systems 28: Annual Conference on Neural Information Processing Systems 2015, December 7-12, 2015, Montreal, Quebec, Canada, pages 3186–3194, 2015.
- [NS12] Praneeth Netrapalli and Sujay Sanghavi. Learning the graph of epidemic cascades. In SIGMET-RICS/Performance, 2012.
- [RD02] Matthew Richardson and Pedro Domingos. Mining knowledge-sharing sites for viral marketing. In KDD, 2002.
- [SS13] Lior Seeman and Yaron Singer. Adaptive seeding in social networks. In 54th Annual IEEE Symposium on Foundations of Computer Science, FOCS 2013, 26-29 October, 2013, Berkeley, CA, USA, pages 459-468, 2013.

A Concentration Bounds

We state the Chernoff bound and Hoeffding's inequality, the two standard concentration bounds used to bound the error of the estimates of the marginal contributions.

Lemma 13 (Chernoff Bound). Let X_1, \ldots, X_n be independent indicator random variables, $X = \sum_{i=1}^n X_i$ and $\mu = \mathbb{E}[X]$. For $0 < \delta < 1$,

$$\Pr[|X - \mu| \ge \delta \mu] \le 2e^{-\mu \delta^2/3}.$$

Lemma 14 (Hoeffding's inequality). Let X_1, \ldots, X_n be independent random variables with values in [0,b]. Let $X = \frac{1}{m} \sum_{i=1}^{m} X_i$ and $\mu = \mathbb{E}[X]$. Then for every $0 < \epsilon < 1$,

$$\Pr[|X - \mathbb{E}[X]| \ge \epsilon] \le 2e^{-2m\epsilon^2/b^2}.$$

B Estimates of Expected Marginal Contributions

Recall that $S \sim \mathcal{D}_a$ is the set obtained from $S \sim \mathcal{D}$ conditioned on $a \in S$. Similarly, we defined $S \sim \mathcal{D}_{-a}$ by conditioning on S not containing a and we also extended this definition for multiple nodes, such as with $\mathcal{D}_{a,-b}$. Similarly as with \mathcal{D}_a and \mathcal{D}_{-a} , let \mathcal{S}_a and \mathcal{S}_{-a} be the collections of samples containing and not containing a respectively. We also extend this notation for multiple nodes, e.g., $\mathcal{S}_{a,-b}$.

The first order marginal contributions v(a) are estimated with the difference $\tilde{v}(a)$ between the average value of samples containing a and the average value of samples not containing a, i.e.

$$\tilde{v}(a) := \left(\frac{1}{|\mathcal{S}_a|} \sum_{S \in \mathcal{S}_a} f(S) - \frac{1}{|\mathcal{S}_{-a}|} \sum_{S \in \mathcal{S}_{-a}} |f(S)|\right).$$

Similarly, the second order marginal contribution $v_b(a)$ are estimated with

$$\tilde{v}_b(a) := \left(\frac{1}{|\mathcal{S}_{a,b}|} \sum_{S \in \mathcal{S}_{a,b}} f(S) - \frac{1}{|\mathcal{S}_{-a,b}|} \sum_{S \in \mathcal{S}_{-a,b}} |f(S)| \right).$$

We show that these estimates $\tilde{v}(a)$ and $\tilde{v}_b(a)$ are arbitrarily good when the distribution \mathcal{D} is a bounded product distribution.

Lemma 15. The estimates $\tilde{v}(a)$ and $\tilde{v}_b(a)$ are arbitrarily close to the first and second order marginal contributions v(a) and $v_b(a)$ of a node a given samples from a bounded product distribution \mathcal{D}^4 . For all $a, b \in N$ and given $\operatorname{poly}(n, 1/\delta, 1/\epsilon)$ i.i.d. samples from \mathcal{D} ,

$$|\tilde{v}(a) - v(a)| \le \epsilon$$
 and $|\tilde{v}_b(a) - v_b(a)| \le \epsilon$

with probability at least $1 - \delta$ for any $\delta > 0$.

⁴A product distribution $S \sim \mathcal{D}$ is bounded if the marginal probabilities are bounded away from 0 and 1, i.e., $1/\operatorname{poly}(n) \leq \Pr[a \in S] \leq 1 - 1/\operatorname{poly}(n)$ for all $a \in V$.

Proof. We give the analysis for the second order marginal contribution. The proof follows identically for the first order marginal contribution by treating b as null. Note that since \mathcal{D} is a product distribution.

$$\underset{S \sim \mathcal{D}: b \in S, a \not \in S}{\mathbb{E}}[f(S \cup a)] - \underset{S \sim \mathcal{D}: b \in S, a \not \in S}{\mathbb{E}}[f(S)] = \underset{S \sim \mathcal{D}: b \in S, a \in S}{\mathbb{E}}[f(S)] - \underset{S \sim \mathcal{D}: b \in S, a \not \in S}{\mathbb{E}}[f(S)].$$

Since marginal probabilities of the product distributions are assumed to be bounded from below and above by $1/\operatorname{poly}(n)$ and $1-1/\operatorname{poly}(n)$ respectively, $|\mathcal{S}_{a,b}| = m/\operatorname{poly}(n)$ and $|\mathcal{S}_{-a,b}| = m/\operatorname{poly}(n)$ for all a by Chernoff bound. In addition, $\max_S f(S)$ is assumed to be bounded by $\operatorname{poly}(n)$. So by Hoeffding's inequality,

$$\Pr\left(\left|\frac{1}{|\mathcal{S}_{a,b}|} \sum_{S \in \mathcal{S}_{a,b}} f(S) - \mathbb{E}_{S \sim \mathcal{D}|b \in S, a \in S}[f(S)]\right| \ge \epsilon/2\right) \le 2e^{-\frac{m\epsilon^2}{\text{poly}(n)}},$$

for $0 < \epsilon < 2$

$$\Pr\left(\left|\frac{1}{|\mathcal{S}_{-a,b}|} \sum_{S \in \mathcal{S}_{-a,b}} |f(S)| - \mathbb{E}_{S \sim \mathcal{D}|b \in S, a \notin S}[f(S)]\right| \ge \epsilon/2\right) \le 2e^{-\frac{m\epsilon^2}{\text{poly}(n)}}$$

for $0 < \epsilon < 2$. Thus,

$$\Pr\left(\left|\underset{S \sim \mathcal{D}|b \in S, a \notin S}{\mathbb{E}} [f(S \cup a) - f(S)] - \left(\frac{1}{|\mathcal{S}_{a,b}|} \sum_{S \in \mathcal{S}_{a,b}} f(S) - \frac{1}{|\mathcal{S}_{-a,b}|} \sum_{S \in \mathcal{S}_{-a,b}} |f(S)|\right)\right| \ge \epsilon\right)$$

$$\leq 4e^{-\frac{me^2}{\text{poly}(n)}}$$

for $0 < \epsilon < 2$.

C Erdős-Rényi Random Graphs

A $G_{n,p}$ Erdős–Rényi graph is a random graph over n vertices where every edge realizes with probability p. Note that the graph obtained by the two step process which consists of first the stochastic block model and then the independent cascade model is a union of $G_{|C|,p_C}$ for each community C. The following are seminal results from Erdős–Rényi characterizes phase transitions for $G_{n,p}$ graphs.

Lemma 3. [ER60] Assume C is a "dense" community, then the subgraph G[C] of G is connected with probability $1 - O(|C|^{-2})$.

Proof. Assume $p_C = c \log |C|/|C|$ for c > 1. From Theorem 4.6 in [BHK] which presents the result from [ER60], the expected number of isolated vertices a in G(|C|, p) is

$$\mathbb{E}[i] = |C|^{1-c} + o(1)$$

and from Theorem 4.15 in [BHK], the expected number of components of size between 2 and |C|/2 is $O(|C|^{1-2c})$. Thus the expected number of components of size at most n/2 is $O(n^{1-c})$ and the probability that the graph is connected is $1 - O(|C|^{1-c})$. Finally, since $c \ge 3$ for dense communities, the probability that the graph for community C is connected is $1 - O(|C|^{-2})$.

Lemma 7. [ER60] Let C be a tight community with $|C| = \omega(1)$, then G[C] has a "giant" connected component containing a constant fraction of the nodes in C w.p. 1 - o(1).

D Missing Analysis from Section 3

Theorem 6. In the dense communities and small seed set setting, COPS with α -overlap allowed, for any constant $\alpha \in (0,1)$ is a 1-o(1)-approximation algorithm for learning to influence from samples from a bounded product distribution \mathcal{D} .

Proof. First, we claim that a node $a \in C$ is not removed from the ordering if there is no other node from C before a. For $b \notin C$, we have

$$v(a) = \mathbb{E}_{S \sim \mathcal{D}_{-a}}[f_S(a)] = \mathbb{E}_{S \sim \mathcal{D}_{-a}}[f_S(a) : b \in S] = \mathbb{E}_{S \sim \mathcal{D}_{-a,b}}[f_S(a)] = v_b(a)$$

where the second equality is since a and b are in different communities and since \mathcal{D} is a product distribution. Thus, OVERLAP (a, b, α) = False for any $\alpha \in (0, 1]$.

Next, recall that $v(a) \leq |C|$ for all $a \in C$. Thus, by Lemmas 4 and 5. COPS returns a set that contains one node from k different communities that have sizes that are at most a factor 1 - o(1) away from the sizes of the k largest communities. Since the k largest communities are connected with high probability, the optimal solution contains one node from each of the k largest communities. Thus, we obtain a 1 - o(1) approximation.

E Missing Analysis from Section 4

Lemma 8. Let C be a loose community, then $f(\{a\}) \leq c$ for all $a \in C$ and some constant c.

Proof. Fix a node $a \in C$. We consider a Galton-Watson branching process starting at individual a where the number of offsprings of an individual is $X = \text{Binomial}(|C| - 1, p_C)$. We show that the expected total size s of this branching process is $1/(1 - p_C \cdot (|C| - 1))$ and that this expected size s upper bounds $f(\{a\})$.

We first argue that $s \geq f(\{a\})$. The expected number of nodes influenced by a can be counted via a breadth first search (BFS) of community C starting at a. The number of edges leaving a node in this BFS is Binomial($|C| - 1, p_C$), which is exactly the number of offsprings of an individual in the branching process. Since the nodes explored in the BFS are only the nodes not yet explored, the number of nodes explored by BFS is upper bounded by the branching process and we get $\mathbb{E}[s] \geq f(\{a\})$.

Next, let $\mu = p_C \cdot (|C| - 1) < 1$ be the expected number of offsprings of an individual in the branching process. Let s_i be the expected number of individuals at generation i of the branching process. We show by induction that $\mathbb{E}[s_i] = \mu^i$. The base case is trivial for i = 1. Next, for i = 2,

$$\mathbb{E}[s_i] = \sum_{j=0}^{\infty} \Pr[s_{i-1} = j] \cdot \mathbb{E}[s_i | s_{i-1} = j] = \sum_{j=0}^{\infty} \Pr[s_{i-1} = j] \cdot j \cdot \mu = \mu \, \mathbb{E}[s_{i-1}] = \mu^i$$

where the last inequality is by the inductive hypothesis. Thus, $\mathbb{E}[s] = \sum_{i=0}^{\infty} \mu^i = 1/(1-\mu)$ since $\mu < 1$. Finally, since C is tight, $\mu \leq 1 - \epsilon$ for some constant $\epsilon > 0$ and $f(\{a\}) \leq \mathbb{E}[s] \leq 1/\epsilon$.

Lemma 11. If $a, b \in C$ and C is a tight community, then $OVERLAP(a, b, \alpha) = True \ for \ \alpha = o(1)$.

Proof. Let $a, b \in C$ s.t. C is a tight community. The marginal contribution of node a_i can be decomposed into the whether $a \in gc(C)$:

$$v(a) = \Pr_G[a \in \gcd(C)] \cdot \underset{S \sim \mathcal{D}_{-a}}{\mathbb{E}} [f_S(a) : a \in \gcd(C)] + \Pr_G[a \not\in \gcd(C)] \cdot \underset{S \sim \mathcal{D}_{-a}}{\mathbb{E}} [f_S(a) : a \not\in \gcd(C)]$$

Since \mathcal{D} is a product distribution,

$$\mathbb{E}_{S \sim \mathcal{D}_{-a}}[f_S(a) : a \in gc(C)] = (\Pr[b \notin gc(C)] + \Pr_{S \sim \mathcal{D}}[b \in gc(C), b \notin S]) \mathbb{E}_{S \sim \mathcal{D}_{-a, -b}}[f_S(a) : a \in gc(C)]$$

$$\geq (1 + \epsilon) \Pr[b \notin gc(C)] \mathbb{E}_{S \sim \mathcal{D}_{-a, -b}}[f_S(a) : a \in gc(C)]$$

$$= (1 + \epsilon) \mathbb{E}_{S \sim \mathcal{D}_{-a, b}}[f_S(a) : a \in gc(C)]$$

for some constant $\epsilon > 0$ since $\Pr_{S \sim \mathcal{D}_{-a}}[b \in gc(C), b \notin S] \ge \epsilon_1$ for some constant $\epsilon_1 > 0$. Since,

$$\Pr_{G}[a \in \gcd(C)] \underset{S \sim \mathcal{D}_{-a}}{\mathbb{E}} \left[f_{S}(a) : a \in \gcd(C) \right] \ge \epsilon' |C| \ge \epsilon' \Pr_{G}[a \notin \gcd(C)] \cdot \underset{S \sim \mathcal{D}_{-a}}{\mathbb{E}} \left[f_{S}(a) : a \notin \gcd(C) \right]$$

for some constant $\epsilon' > 0$, we get

$$v(a) \ge (1+\epsilon) \Pr_{G}[a \in \gcd(C)] \cdot \underset{S \sim \mathcal{D}_{-a,b}}{\mathbb{E}} [f_{S}(a) : a \in \gcd(C)]$$

$$+ \Pr_{G}[a \notin \gcd(C)] \cdot \underset{S \sim \mathcal{D}_{-a}}{\mathbb{E}} [f_{S}(a) : a \notin \gcd(C)]$$

$$\ge (1+\epsilon\epsilon'/2) \Big(\underset{G}{\Pr}[a \in \gcd(C)] \underset{S \sim \mathcal{D}_{-a,b}}{\mathbb{E}} [f_{S}(a) : a \in \gcd(C)]$$

$$+ \underset{G}{\Pr}[a \notin \gcd(C)] \cdot \underset{S \sim \mathcal{D}_{-a,b}}{\mathbb{E}} [f_{S}(a) : a \notin \gcd(C)] \Big)$$

$$= (1+\epsilon\epsilon'/2) v_{b}(a)$$

Thus, OVERLAP (a, b, α) = True for $\alpha = o(1)$.

Theorem 12. With overlap allowed $\alpha = 1/\operatorname{poly}(n)$, COPS is a constant factor approximation algorithm for learning to influence from samples drawn from a bounded product distribution \mathcal{D} in the setting with tight and loose communities and non-ubiquitous seed sets.

Proof. First, observe that $f(S) \geq k$ since nodes trivially influence themselves. Let a_i be the node picked by the algorithm that is in the ith position of the ordering after the pruning and assume $i \leq t$. By Lemma 10, $f(\{a_i\}) \geq \epsilon |C_i|$ where C_i is the ith largest tight community with super-constant individual influence and with at least one node in S^* . Thus $a_i, i \in [t]$ is in a tight community, otherwise it would have constant influence by Lemma 8, which is a contradiction with $f(\{a_i\}) < \geq \epsilon |C_i|$. Since $a_i, i \in [t]$ is in a tight community, by Lemma 11, we obtain that a_1, \ldots, a_i are all in different communities. We denote by S_t the subset of the solution returned by COPS and obtain We obtain $f(S^*) \leq \sum_{i=1}^t |C_i| + c \cdot k \leq \sum_{i=1}^t \frac{1}{\epsilon} \cdot f(\{a_i\}) + c \cdot f(S) = \frac{1}{\epsilon} \cdot f(S_t) + c \cdot f(S) \leq c_1 \cdot f(S)$ for some constant ϵ, c, c_1 by Lemmas 9, 10, and since a_i, a_j are in different communities for $i, j \leq t$. \square

F Additional Experimental Setup and Analysis

Additional description for the experimental setup. Each point in a plot corresponds to the average performance of the algorithms over 10 trials. The default values for k is k = 10. For the experiments on synthetic data, the default overlap allowed is $\alpha = 0.5$, for the Facebook experiments $\alpha = 0.4$ and for the DBLP experiments $\alpha = 0.2$. The default edge weights are chosen so that in the random realization of G the average degree of the nodes is approximately 1.

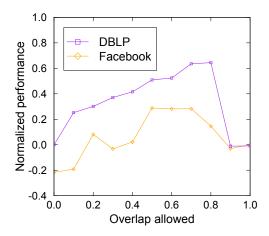


Figure 3: Performance of COPS as a function of the overlap α allowed. The performance is normalized so that the performance of GREEDY and MARGI corresponds to value 1 and 0 respectively

Additional analysis. As discussed in Section 5, the performance of COPS as a function of the overlap allowed (Figure 3) can be explained as follows: Its performance slowly increases as the the overlap allowed increases and COPS can pick from a larger collection of nodes until it drops when it allows too much overlap and picks mostly very close nodes from the same largest community.